

Colorful Image Colorization with TensorFlow

Sophia Schulze-Weddige

Malin Spaniol

Maren Born

Implementing Artificial Neural Networks with Tensorflow

Universität Osnabrück

April 19, 2020

Contents

1	Introduction	3
2	Theoretical Background	3
2.1	Related work	3
2.2	Convolutional neural networks for image colorization	4
2.3	Guiding Paper	5
3	Network Structure and Implementation	6
3.1	Data Generation	6
3.2	Model Structure	8
3.3	Training	9
3.4	Testing	10
4	Results	10
4.1	Classical Approach	10
4.2	Classification Approach	12
5	Discussion	13
6	Conclusion	14
7	Literature	15
A	Error Message	16
B	Layers of the Model	17

1 Introduction

Based on the paper *Colorful Image Colorization* (Zhang et al., 2016), this project aims to reimplement a similar artificial neural network which transforms grayscale images into colorful pictures. This task involves creating a dataset based on pictures that are converted into the CIELAB colorspace (Lab), such that the lightness channel L can be considered as the input whereas the a and b channels, which encode color information, form the target for the model. We closely rebuild the layers of the original model (they used “caffe” (Jia et al., 2014)) using tensorflow 2.0. Our project is divided into two steps, which are also discussed in the paper (Zhang et al., 2016). In the first step, the unaltered ab channels are utilized as the prediction target, applying the mean squared error loss. As predicted by Zhang et al. (2016) this approach favors desaturated colors predicting images to appear sepia or greyish. This approach can be thought of as the “classical” way of automated image colorization with convolutional neural networks.

The second step aims to rebuild the main contribution to the image colorization problem from Zhang et al. (2016), namely to translate the problem to a classification task. By doing so, Zhang et al. (2016) were able to predict more plausible colors for the grayscale images. Due to the lack of time, we were not able to train our models long enough to reproduce the good results provided. More training time will hopefully lead to vivid and realistic colorizations for gray scale images with our model as well.

2 Theoretical Background

Image colorization can be used to modernize pictures or movies. There is a wide range of methods, ranging from hand colorization with photoshop to automatic colorization with artificial intelligence. For examples on images colorization see <https://www.reddit.com/r/Colorization/>. The TV-series “Greatest Events of WWII in Colour” states a good example for the usage of colorized footage. They use the techniques such that events appear more contemporary thus the historic incidents turn out even more worrying (<https://www.imdb.com/title/tt9103932/>). Automating the colorization problem with deep learning seems to be achievable, as training data is easily available.

2.1 Related work

Automatic approaches solving the colorization problem mostly differ in acquisition and handling of the data in order to model the accurate correspondence (Zhang et al., 2016). One can differentiate between parametric and non-parametric approaches. Non-parametric methods predict colors based on one or more reference images. That means the color distribution of the refer-

ence images, which are either provided by the user (e.g. Scribble-based colorization by Levin et al. (2004)) or automatically, is transferred to the target image. Hence, performance depends heavily on the quality of the provided data (Cheng et al., 2015).

Parametric methods on the other hands, learn prediction functions from large datasets of color images. Different methods are available to achieve this, one of which are convolutional neural networks (CNNs) in which the problem can be posed as regression or classification of quantized color values (Zhang et al., 2016).

2.2 Convolutional neural networks for image colorization

Whenever dealing with image data, CNNs are frequently used as model structures for deep learning tasks. They are inspired by the visual cortex of the brain. The idea is that highly specialized components learn a very specific task, which is similar to the receptive fields of neurons in the visual cortex (Hubel and Wiesel, 1962). These components can be combined to high-level features, which again can be merged to classes or transformed to the desired output shape. In CNNs, this concept is implemented by several successive convolutional layers: a weight kernel moves over the input image and calculates the new pixel value for each pixel position by multiplying the weights of the kernel with the neighboring pixel values and summing them up (see fig. 1). Different kernels can generate different so-called feature maps. One feature map targets the same feature (e.g. edges) in different image sections.

In this manner, CNNs can store spatial information about pixels and features. In a subsequent pooling layer dimensions are reduced by summarizing over the image section (e.g. max pooling takes the highest value of a certain image section). This facilitates the computation and drops unnecessary information (Effenberger, 2019). In recent years, CNNs improved such that they outperform humans in many classification tasks (Russakovsky et al., 2014).




Figure 1: Convolutional operation: kernel slides over the input, multiplying it with its weight before summarizing the 3x3 neighborhood (Escontrela, 2018).

2.3 Guiding Paper

Zhang and colleagues (2016) propose a fully automatic approach to colorize grayscale images. They choose to solve this task with a feed-forward CNN as a classification task using a custom loss and class-rebalancing at training time in order to increase the diversity of the colours in the final results. The model is trained over a million color images from ImageNet dataset (Rusakovsky et al., 2014). Concerning the CNN architecture, they use a single stream, VGG-styled network with added depth and dilated convolutions.

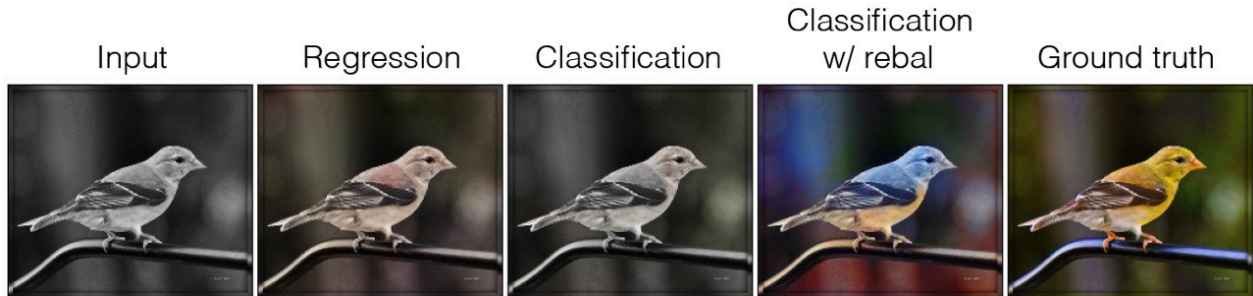


Figure 2: Example results from different approaches tested by Zhang et al. (2016)

The mapping, which Zhang and colleagues (2016) aim to learn, results out of the *Lab* color space: The input is the *L* lightness channel, the target channels are the *a* (from green to red) and *b* (from blue to yellow) color channels. First, they used the Euclidean loss to train the model. This leads to overall grayish colors because the Euclidean loss favors the mean of all color pixel values. The example results in figure 2 show this phenomenon in the column “regression”. Therefore, Zhang et al. (2016) choose to treat the problem as a multinomial classification task. The *ab* output space is therefore divided into bins of grid size ten, and the 313 color values in-gamut span the 313 possible color combinations. Following, a mapping to a probability distribution over all 313 possible colors is learned for each input. The ground truth color is converted to a vector, using a soft-encoding scheme and a multinomial cross entropy loss, which is responsible for the class-rebalancing. Thereafter, they map the probability distribution to the color values. The class-rebalancing operates pixel-wise and the loss of each pixel is re-weighted at training time, based on how often the color occurs.

3 Network Structure and Implementation

The re-implementation of the original paper (Zhang et al., 2016) was conducted in two steps and can be found on GitHub (https://github.com/marumse/colorize_images). This chapter will explain these two crucial steps in further detail. In both approaches, the ImageNet2012 dataset was used. Firstly, we implemented the model structure as described in the paper (see figure 3) and trained the model with the color layers of the input images as the target. Secondly, we translated the colorization task to a classification problem and trained the same model structure with the altered problem representation.




Figure 3: The network architecture of Zhang et al. (2016).

3.1 Data Generation

The ImageNet2012 dataset was available to us through the university server. As we were working with a large amount of data in each batch, it would have been impossible to load the whole dataset at once, hence we used a data generator to load the input and target images successively for each batch. Although there are inbuilt data generators available from keras that allow for some means of data augmentation, we built our costum generator to ensure the functionality we were aiming at. The generator takes the batch size and a list which contains the paths to the images that should be used to create the input and target arrays as described in the following. First, the images are loaded and resized to a uniform shape of (224, 224, 3), which corresponds to the height, the width and the number of the color channels, respectively. Then, the images are transformed from BGR to *Lab* color space. The lightness channel (*L*-layer), which displays most of the structure, is separated from the other two channels and used as the input for the model. The remaining two channels (*ab*-layers) encode the color information of the images and are used as the models' target in the first approach. In the second approach, these processing steps remain the same. Additionally the target arrays are transformed to become a classification task.

```

95 def generate_data(batch_size, file_list):
96     """ Replaces Keras' native ImageDataGenerator.
97         This function is a data generator that loads a customized version of some data. More precisely, it loads images,
98         tranforms them into LAB color space and returns the first layer as the input and the other two layers as the target for the model
99     Args:         batch_size
100                  file_list containing all image paths
101     Return:      a tuple containing a numpy array with the inputs and a second numpy array with the targets
102     """
103     i = 0
104     while True:
105         image_batch = []
106         label_batch = []
107
108         for b in range(batch_size):
109             # shuffle data when all files have been seen
110             if i == len(file_list):
111                 i = 0
112                 np.random.shuffle(file_list)
113             sample = file_list[i]
114             i += 1
115             image = cv2.resize(cv2.imread(sample), (224,224))
116             image = cv2.cvtColor(image, cv2.COLOR_BGR2LAB)
117             # split the image into the L layer for the input and ab layers for the target
118             L = image[:, :, 0][ :, :, np.newaxis]
119             ab = image[:, :, 1:]
120             # reduce the number of colors to 121 (11 different a and b values respectively)
121             ab = (ab//25)*25
122             # use cantor pairing to combine layer a and b
123             cantor = cantor_pairing(ab)
124             # make a one-hot-encoding out of the cantor labels
125             hot = one_hot_encoding(cantor)
126             # append both to the corresponding lists
127             image_batch.append(L)
128             label_batch.append(hot)
129         yield (np.array(image_batch), np.array(label_batch))

```

Figure 4: The data generator creates batches of input and target tuples.

The transformation of images for the classification task is done in three main steps. The first step is to discretize the continuous color space and reduce the number of possible colors by quantizing the a and b color ranges into eleven bins each. This yields a total of 121 possible colors by combining the a - and b -layers. In the second step, the a - and b -layers are combined into a single layer, which keeps the same height and width dimensions as before. This means that the color information of each pixel is now encoded in a single number rather than two, which leads to a single target layer. Cantor pairing is used to generate a unique and deterministic number from the two a and b values of each pixel. The Cantor pairing formula is shown in (1) and its implementation in figure 5. As cantor pairing is reversible, one can easily translate the pairing result back to the original color values with no loss of information (Meri, 2007).

$$z = \pi(x, y) = \frac{(x + y + 1)(x + y)}{2} + y \quad (1)$$

```

73 def cantor_pairing(ab):
74     """ Cantor pairing calculates a deterministic and unique number for a pair of positive integers.
75     Args:      two arrays of integers in our case the a and b color channels
76     Return:    the corresponding array with unique values
77     """
78     a = ab[:, :, 0].astype(np.uint32)
79     b = ab[:, :, 1].astype(np.uint32)
80     c = ((a + b) * (a + b + 1)) / 2 + b
81     return c

```

Figure 5: Implementation of the Cantor Pairing Function.

Now that there is a single value encoding each pixels' color, the third and last step is to translate this value into a one-hot encoding. With the help of a dictionary, the cantor values are translated to numbers from 0 to 120. These numbers serve as the index in the one-hot encoding. Hence, color values that were previously represented in two values (a and b color channels) are now encoded by their index in the one-hot encoding. The target array has a shape of (224, 224, 121) and a one-hot vector is located at each pixel position.

```

51 def one_hot_encoding(cantor):
52     """ Get one-hot-encoding for cantor target image.
53     Args:      cantor transformed ab layers from input images
54     Return:    one-hot-encoding with shape width x height x number of colors (224 x 224 x 121)
55     """
56     one_hot = np.zeros((cantor.shape[0], cantor.shape[1], 121))
57     for i, unique_value in enumerate(np.unique(cantor)):
58         index = pair_to_index[unique_value]
59         one_hot[:, :, index][cantor == unique_value] = 1
60     return one_hot

```

Figure 6: Implementation of the one-hot encoding.

3.2 Model Structure

The model structure is equivalent to the original implementation by Zhang et al. (2016). Only minor changes occur, mostly due to the translation from caffe to keras. The model consists of eight blocks comprising two or three repeated convolution and ReLU activation layers followed by a batch normalization. There are no pooling layers in the model, as changes in resolution are achieved through spatial down- or upsampling between the convolution blocks. A transposed convolutional layer is used to inverse the convolution and upsample the output back to the correct size. The model summary can be looked at in the appendix B (figure 12).

In the first approach, a stochastic gradient descent optimizer (SGD) with a learning rate of 0.001 and a momentum of 0.9 is used, which is inspired by the original paper. SGD is a stochastic approximation to gradient descent, which uses an estimation for the gradients and

thereby decreases the computational complexity. In each optimization step, a random data point is selected from which the gradients are calculated rather than using the whole dataset. The downside to SGD is that it does not guarantee to converge to a solution as the gradients might vary heavily from sample to sample (Effenberger, 2019). Further, mean squared error (MSE) is implemented as the loss function, which uses the average squared difference between the estimated and the target values (2). In the case of image data, that means that each estimated pixel value is compared to the target pixel value and the average over the squared differences is used to optimize the estimation process.

$$MSE = \frac{1}{n} \sum_{i=1}^N (y_i - \tilde{y}_i)^2 \quad (2)$$

In the second approach, the softmax activation function is used in the last layer to prepare the model’s output for the categorical cross entropy loss (3), which compares the output vector (\tilde{y}) to the one-hot encoded target vector (y).

$$L(y, \tilde{y}) = - \sum_{j=0}^M \sum_{i=0}^N y_{ij} * \log(\tilde{y}_{ij}) \quad (3)$$

Both, SGD and Adam optimizer are tested. In the two approaches, kernel weights are initialized with the *Glorot uniform initializer* which draws samples from a uniform distribution depending on the number of input and output units of the corresponding weight tensor (Glorot and Bengio, 2010). Biases are initialized with zeros.

3.3 Training

The models were trained on the grid of the Institute of Cognitive Science at Osnabrueck University. A helper script was written that distributed several grid jobs on different computers in order to experiment with the hyperparameters such as the learning rate and the batch size. Through this script, one can easily change these parameters as well as select the dataset in addition to the environment and switch between the training mode and the prediction mode. The classical model was trained with 5000 images, a batch size of 10 and a learning rate of 0.001. The classification model was trained with 2000 images, batch sizes of 10 or 20 and learning rated from 0.1 to 0.001. Unfortunately, the grid jobs terminated after three to six epochs without giving a hint for errors in the code. It appears to be a grid internal memory problem (see appendix A, figure 11). Therefore, we saved the model in checkpoints after each epoch and used those weights for the testing.

3.4 Testing

To evaluate the predictive power of our two approaches, colors for unseen test images are predicted and evaluated via visual inspection. When the script is started in prediction mode, the model is not trained, but the weights from the corresponding training process are loaded. The model then predicts the most plausible colors for the L -layer of each input image.

In the first approach, the model’s output can be interpreted as the ab layers. That means the output can be combined with the input (L -layer) and displayed as an image straight away.

For the second approach, three decoding steps are necessary before generating human-readable images. Firstly, the output of the softmax layer is decoded to find the index of the most likely color value. Secondly, this index is translated to the cantor pairing value it represents with the help of the created dictionary. And thirdly, the cantor pairing value is transformed back to the a and b values it is originally constituted of. These a and b layers can finally be combined with the L -layer to display the predicted image in Lab color space.

4 Results

In the following section we will present the results of our classical approach and our classification approach separately. As the task of image colorization can be best evaluated by visual inspection we provide example images that were generated with the two different approaches.

4.1 Classical Approach

When looking at the predicted images, one can see that the structure of the images is still intact and it can be easily recognized what is depicted in the image. This is due to the fact that in Lab color space the L -layer displays most of the structure. The model weights with which these images were generated predict monotonous color values, with a value of 131 for the a channel and 136 for the b channel. These two values are assigned to all pixels throughout all images and correspond to a grayish sepia hue. There are neither differences in color between pixels nor images.

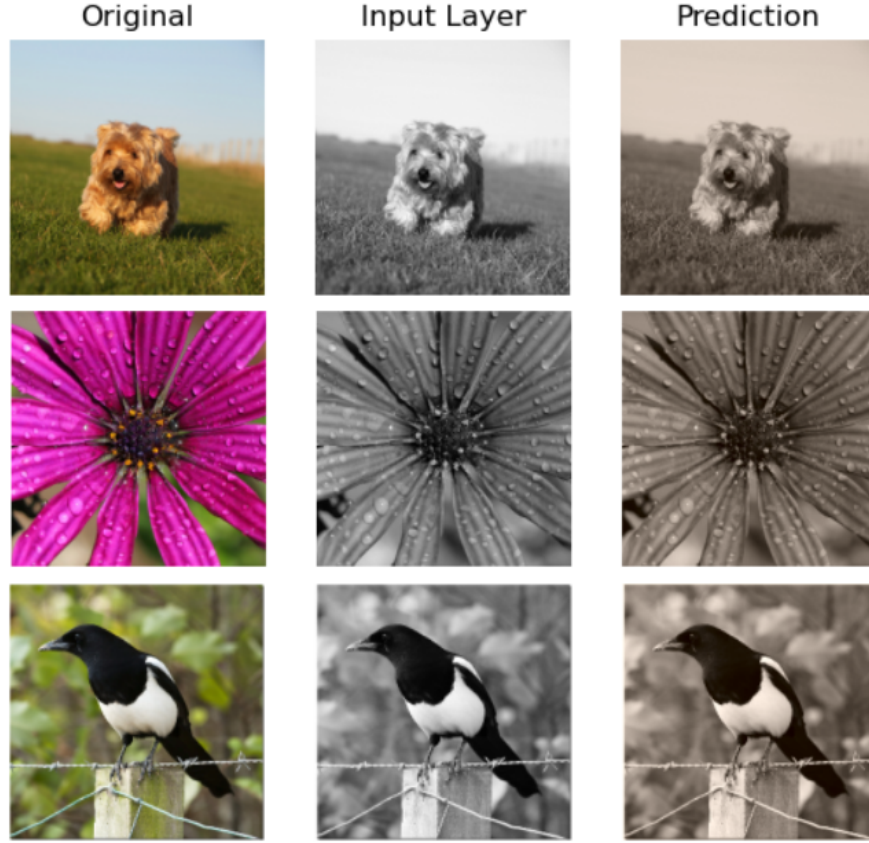


Figure 7: Example ground truth images from the ImageNet test set, with corresponding extracted L -input layer and prediction by our classical approach.

As shown in figure 8, both the training and the validation loss start with a very high value of around 17000 for the training loss and 13000 for the validation loss. Within the first epochs, they decrease rapidly and then converge to a value around 230 to 245. In the beginning, the validation loss is slightly lower than the training loss.




Figure 8: History of the MSE loss: The loss is evaluated after each epoch, for a period of 50 epochs.

4.2 Classification Approach

The first thing one might notice is that the originals look different in this approach. This is due to the fact that the color space was quantized to 121 different color values. This means that all color values are distributed into bins and, hence, no gradual changes in color can be observed in the images. The predictions display a similar grayish color hue throughout the image as in the classical approach. The model weights with which these images were generated predict an a and b value of 125 for most pixels, but additionally, some pixels in the b channel have a predicted value of 100. Those differently predicted pixels appear systematically in a grid of 14×14 as little dots throughout each image. Each of those dots consists of nine pixels with the value 100. For the a channel a uniform value of 125 is predicted for all pixels.

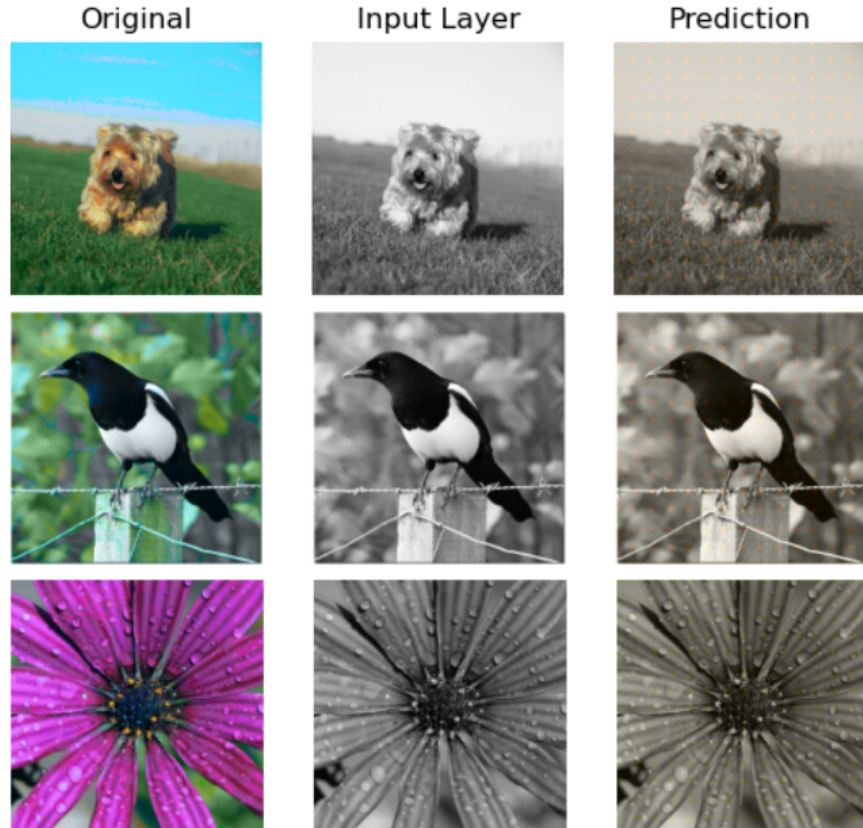


Figure 9: Example ground truth images from the ImageNet test set, with corresponding extracted L -input layer and prediction by our classification approach.

5 Discussion

As predicted by Zhang et al. (2016), using a regression loss function like the MSE as in the first approach, leads to desaturated colors. This indicates that the model learns the mean color of the image dataset and predicts that mean value for all pixels. This is due to the nature of Euclidean loss and MSE, which is not robust to the multimodal nature of the colorization problem (Zhang et al., 2016). The optimal solution to such loss functions will always be the mean of the whole color set.

The rapidly decreasing loss (see figure 8) suggests that the learning process of the classical model happens quickly and converges after the mean value is reached. With a longer training period and more training images different color variations in the predicted images can be expected though still desaturated (Zhang et al., 2016).

In order to improve the quality of the results, the MSE loss can be replaced by a custom loss function. In our classification approach, we implemented a categorical cross entropy loss, to improve the first approach. The transformation to a classification problem promises vivid and realistic colors, which we were not able to reproduce. The short training periods also pose a problem for this implementation. Additionally, an unexpected bug occurred, that predicted colored dots on the output images rather than coloring the whole image (see figure 10). The grid jobs always terminated after three to six epochs when training the classification approach. Thus, the predictions had to be made with the intermediate checkpoints that were saved after each improving epoch. Therefore, it is not guaranteed that this surprising prediction of dots remains if training for a longer period of time is possible. As the dots appear on an evenly distributed grid of 14 by 14, we suspect that the error might occur in the transposed convolutional layer. To be more precise, in the last convolutional block the data is represented with 14 by 14 neurons and the transposed convolution upsamples this output to the desired size of 224 by 224 neurons. We thoroughly tested the helper functions for encoding the color channels of the images as one-hot vectors, and for decoding them back to images. Therefore, we assume that the artifact results from a conversion problem at some point, but as no error message occurs, and the overall process is a complex calculation, we were not able to find the cause.

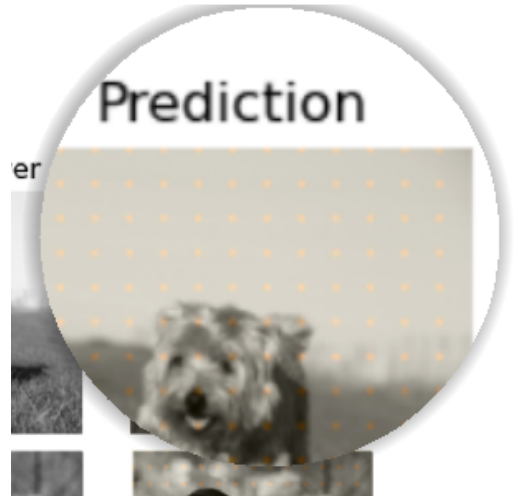


Figure 10: Closeup to the colored dots in the prediction of the classification model.

The first improving steps to consider would be to increase the number of training images, the batch size and the number of epochs. We are aware that the values we chose for those three parameters are too small to expect an excellent result. As computation turned out to be very time consuming, even when performed on the grid of the Institute of Cognitive Science, we were not able to run the model for more episodes with more time consuming but promising parameters (e.g. higher batch size).

A next step to improve the model would be to implement a class-rebalancing by weighing color bins, depending on how frequently they occur through the whole dataset. Pixels of rare color bins should be weighted more. This could be a way to yield more vivid and realistic color results. Translating a colorization problem into a classification problem is a new approach. Therefore, not many solutions or partial solutions existed for implementation steps, such as an adequate data generator. Although the code from the guiding paper was available to us in caffe, the translation of the problem was not deductible from their documentation. Hence, we came up with our own ideas for the implementation resulting in the data generator with the cantor pairing and one-hot encoding as it can be found in figure 4-6.

6 Conclusion

In this project, we implemented a convolutional neural network predicting color values for grayscale images (Zhang et al., 2016) and trained it on the ImageNet2012 dataset. The results of our first step resemble the findings of the paper, where they used the Euclidean loss (Zhang et al., 2016). The predicted images appear to be sepia, as expected. The results of our second step could not replicate the same color quality as demonstrated in the paper. Even though we converted the task to a classification problem in a similar way and implemented our own quantization, we did not succeed to reach colorful results. Larger training datasets and batch sizes might improve the results, but primarily further investigations are needed to explain the unexpected prediction of tiny colored dots throughout the output images.

7 Literature

- Cheng, Z., Yang, Q., and Sheng, B. (2015). Deep colorization. *2015 IEEE International Conference on Computer Vision (ICCV)*.
- Effenberger, L. (2019). Implementing anns with tensorflow.
- Escontrela, A. (2018). Towards data science: Convolutional neural networks from the ground up. <https://towardsdatascience.com/convolutional-neural-networks-from-the-ground-up-c67bb41454e1>.
- Glorot, X. and Bengio, Y. (2010). Understanding the difficulty of training deep feedforward neural networks. In Teh, Y. W. and Titterton, D. M., editors, *AISTATS*, volume 9 of *JMLR Proceedings*, pages 249–256. JMLR.org.
- Hubel, D. H. and Wiesel, T. N. (1962). Receptive fields, binocular interaction and functional architecture in the cat's visual cortex. *The Journal of Physiology*, 160(1):106–154.
- Jia, Y., Shelhamer, E., Donahue, J., Karayev, S., Long, J., Girshick, R., Guadarrama, S., and Darrell, T. (2014). Caffe: Convolutional architecture for fast feature embedding. *arXiv preprint arXiv:1408.5093*.
- Levin, A., Lischinski, D., and Weiss, Y. (2004). Colorization using optimization. *ACM Trans. Graph.*, 23(3):689–694.
- Meri, L. (2007). Some remarks on the cantor pairing function some remarks on the cantor pairing function. *Le Matematiche*, 62.
- Russakovsky, O., Deng, J., Su, H., Krause, J., Satheesh, S., Ma, S., Huang, Z., Karpathy, A., Khosla, A., Bernstein, M. S., Berg, A. C., and Li, F. (2014). Imagenet large scale visual recognition challenge. *CoRR*, abs/1409.0575.
- Zhang, R., Isola, P., and Efros, A. A. (2016). Colorful image colorization. *CoRR*, abs/1603.08511.

A Error Message

```
[var/lib/gridengine/util/starter.sh: line 41: 21842 Killed                  /usr/bin/cgexec -g freezer,memory,cpuset:${C
GPATH} $@

Lavish subscription of memory-resources:

    32768 Mbytes requested,
    but peak usage was only 16877.6 Mbytes.

-----> Please request grid resources more carefully! <-----
```

Figure 11: Error message that occurs during running the classification model on the university server.

B Layers of the Model

Layer (type)	Output Shape	Param #	conv2d_13 (Conv2D)	(None, 28, 28, 512)	2359808
conv2d_1 (Conv2D)	(None, 224, 224, 64)	640	activation_13 (Activation)	(None, 28, 28, 512)	0
activation_1 (Activation)	(None, 224, 224, 64)	0	batch_normalization_5 (Batch Normalization)	(None, 28, 28, 512)	2048
conv2d_2 (Conv2D)	(None, 112, 112, 64)	36928	conv2d_14 (Conv2D)	(None, 28, 28, 512)	2359808
activation_2 (Activation)	(None, 112, 112, 64)	0	activation_14 (Activation)	(None, 28, 28, 512)	0
batch_normalization_1 (Batch Normalization)	(None, 112, 112, 64)	256	conv2d_15 (Conv2D)	(None, 28, 28, 512)	2359808
conv2d_3 (Conv2D)	(None, 112, 112, 128)	73856	activation_15 (Activation)	(None, 28, 28, 512)	0
activation_3 (Activation)	(None, 112, 112, 128)	0	conv2d_16 (Conv2D)	(None, 28, 28, 512)	2359808
conv2d_4 (Conv2D)	(None, 56, 56, 128)	147584	activation_16 (Activation)	(None, 28, 28, 512)	0
activation_4 (Activation)	(None, 56, 56, 128)	0	batch_normalization_6 (Batch Normalization)	(None, 28, 28, 512)	2048
batch_normalization_2 (Batch Normalization)	(None, 56, 56, 128)	512	conv2d_17 (Conv2D)	(None, 28, 28, 512)	2359808
conv2d_5 (Conv2D)	(None, 56, 56, 256)	295168	activation_17 (Activation)	(None, 28, 28, 512)	0
activation_5 (Activation)	(None, 56, 56, 256)	0	conv2d_18 (Conv2D)	(None, 28, 28, 512)	2359808
conv2d_6 (Conv2D)	(None, 56, 56, 256)	590080	activation_18 (Activation)	(None, 28, 28, 512)	0
activation_6 (Activation)	(None, 56, 56, 256)	0	conv2d_19 (Conv2D)	(None, 28, 28, 512)	2359808
conv2d_7 (Conv2D)	(None, 28, 28, 256)	590080	activation_19 (Activation)	(None, 28, 28, 512)	0
activation_7 (Activation)	(None, 28, 28, 256)	0	batch_normalization_7 (Batch Normalization)	(None, 28, 28, 512)	2048
batch_normalization_3 (Batch Normalization)	(None, 28, 28, 256)	1024	conv2d_20 (Conv2D)	(None, 14, 14, 256)	1179904
conv2d_8 (Conv2D)	(None, 28, 28, 512)	1180160	activation_20 (Activation)	(None, 14, 14, 256)	0
activation_8 (Activation)	(None, 28, 28, 512)	0	conv2d_21 (Conv2D)	(None, 14, 14, 256)	590080
conv2d_9 (Conv2D)	(None, 28, 28, 512)	2359808	activation_21 (Activation)	(None, 14, 14, 256)	0
activation_9 (Activation)	(None, 28, 28, 512)	0	conv2d_22 (Conv2D)	(None, 14, 14, 256)	590080
conv2d_10 (Conv2D)	(None, 28, 28, 512)	2359808	activation_22 (Activation)	(None, 14, 14, 256)	0
activation_10 (Activation)	(None, 28, 28, 512)	0	conv2d_23 (Conv2D)	(None, 14, 14, 121)	31097
batch_normalization_4 (Batch Normalization)	(None, 28, 28, 512)	2048	activation_23 (Activation)	(None, 14, 14, 121)	0
conv2d_11 (Conv2D)	(None, 28, 28, 512)	2359808	conv2d_transpose_1 (Conv2DTr)	(None, 224, 224, 121)	131890
activation_11 (Activation)	(None, 28, 28, 512)	0	conv2d_24 (Conv2D)	(None, 224, 224, 121)	14762
conv2d_12 (Conv2D)	(None, 28, 28, 512)	2359808	Total params: 31,420,181		
activation_12 (Activation)	(None, 28, 28, 512)	0	Trainable params: 31,415,189		
			Non-trainable params: 4,992		

Figure 12: Model summary of the implemented CNN, giving detailed information about the different layers, the output shape and the number of parameters.